

DOCUMENT 16

Developer Guide

Contributing, project structure, adding patterns, and the test suite.

Who This Guide Is For

This guide is for engineers working on the NGW core engine, frontend, or CLI tooling. It assumes familiarity with Python 3.11+, FastAPI, and React 18. It covers: project layout, backend architecture, how to add or modify a lighting pattern, the test and benchmark systems, and the contribution workflow.

Python 3.11

Backend runtime

FastAPI

REST API layer

React 18

Frontend SPA

Pytest

Test framework

Project Structure

ngw-core — top-level directory layout

```
ngw-core/
├── engine/                                # Core analysis pipeline
│   ├── orchestrator.py                  # analyze_image() entry point
│   ├── vision_pipeline.py               # Image processing + MediaPipe
│   ├── vlm_adapter.py                   # AI Gateway VLM calls
│   ├── solver.py                        # Weighted consensus solver
│   ├── blueprints/                      # Pattern YAML definitions
│   ├── learning/                        # Failure detection + candidates
│   └── services/                        # shoot_match_service.py etc.
├── api/
│   ├── main.py                          # FastAPI app + lifespan
│   └── routes/                          # shoot_match.py, lab.py, signals.py ...
├── db/
│   ├── database.py                      # init_db(), engine, session
│   └── models.py                        # SQLAlchemy table models
├── ui/
│   ├── src/
│   │   ├── components/                 # React components
│   │   ├── pages/                     # Route-level pages
│   │   └── hooks/                      # Custom React hooks
│   └── vite.config.ts
├── tests/                               # pytest test suite
├── scripts/                             # CLI tools (see Lab Guide)
├── docs/                                # Screenshots, PDFs, markdown
└── data/                               # SQLite DB (dev only)
```

Backend Architecture

The FastAPI backend is structured around three layers: routes (HTTP handlers), services (orchestration logic), and the engine (analysis pipeline). Routes are thin — they validate input, call services, and format responses. Services own business logic. The engine is stateless and pure.

Layer	File(s)	Responsibility
Routes	api/routes/*.py	HTTP validation, auth checks, response formatting
Services	engine/services/*.py	Orchestrate engine calls, build response models
Engine	engine/orchestrator.py	analyze_image() — calls pipeline solver blueprint
Pipeline	engine/vision_pipeline.py	Image decode, MediaPipe, region extraction
VLM	engine/vlm_adapter.py	AI Gateway calls, multi-pass, parse responses
Solver	engine/solver.py	Weighted consensus across VLM passes
Blueprints	engine/blueprints/*.yaml	Pattern definitions: tiers, thresholds, copy
DB	db/database.py + db/models.py	SQLAlchemy sessions, table models

Adding a New Lighting Pattern

Patterns are defined in YAML files under `engine/blueprints/`. The engine loads all YAML files at startup. Adding a pattern requires three steps: write the YAML definition, add gold set images, run benchmarks.

engine/blueprints/loop.yaml — example pattern definition

```
pattern_id:    loop
display_name:  Loop
description:    "Key 45° off-axis, slight elevation. Nose-shadow loops down."

detection:
  confidence_threshold: 0.72
  primary_cues:
    - catchlight_position: [45, 315]    # degrees off center
    - shadow_direction:    loop          # named shadow type
    - fill_ratio_range:    [1.5, 4.0]    # key:fill stop range

blueprint:
  tiers:
    - tier: 1    # exact match
      key:    { type: monolight, modifier: octa_90cm, angle: 45, height: 7.5 }
      fill:    { type: v_flat, side: camera_left }
    - tier: 2    # close substitute
      key:    { type: speedlight, modifier: umbrella_60cm, angle: 45, height: 6.5 }
      fill:    { type: reflector, side: camera_left }

shoot_mode:
  step_order: [camera, key, fill, test, evaluate]
  ceiling_min_ft: 8
```

REQUIRED BLUEPRINT FIELDS

- `pattern_id` — must be unique, lowercase, underscore-separated.
- `confidence_threshold` — set conservatively (0.65–0.80). Too high = misses; too low = false positives.
- At least one tier-1 blueprint entry. Tier-5 (ambient) is optional but recommended.
- `shoot_mode.ceiling_min_ft` — prevents illegal placements on low ceilings.

Write the YAML

1

Create `engine/blueprints/{pattern_id}.yaml` following the schema above.

Validate YAML

2

Run: `python3 scripts/validate_system_yamls.py --strict`

Add Gold Set images

3

Add 3–5 labeled reference images via the Lab `Gold Sets` + `Add New`.

Run benchmarks

4

Run: `python3 scripts/run_benchmarks.py`. New pattern shows as `NOT_TESTED`.

Tune thresholds

5

Adjust `confidence_threshold` until `PASS` rate $\geq 75\%$ on gold sets.

Write tests**6**

Add test cases to `tests/test_engine.py` covering detection and blueprint generation.

Submit PR**7**

Include: YAML, test results screenshot, benchmark output, gold set IDs.

Adding a New API Endpoint

All API endpoints are FastAPI route functions in `api/routes/`. Follow the existing patterns: validate with Pydantic, call a service function, return a typed response model.

`api/routes/example.py` — minimal endpoint pattern

```
from fastapi import APIRouter, Depends, HTTPException
from pydantic import BaseModel
from db.database import get_db
from sqlalchemy.orm import Session

router = APIRouter(prefix="/api/example", tags=["example"])

class ExampleRequest(BaseModel):
    name: str
    value: float

class ExampleResponse(BaseModel):
    result: str
    processed: bool = True

@router.post("/", response_model=ExampleResponse)
async def create_example(
    body: ExampleRequest,
    db: Session = Depends(get_db)
) -> ExampleResponse:
    # Call service layer — never put business logic in routes
    result = example_service.process(db, body.name, body.value)
    return ExampleResponse(result=result)
```


api/main.py — register the router



```
# In api/main.py, import and include the new router:
from api.routes.example import router as example_router

app.include_router(example_router)
# Endpoint is now live at /api/example/
# Visible in Swagger UI at http://localhost:8000/docs
```

Convention	Rule
Route prefix	All routes use /api/ prefix. Lab routes use /api/lab/.
Auth	Lab routes require x-lab-email header checked against LAB_ALLOWED_EMAILS.
Error responses	Raise HTTPException with status_code and detail. Never return raw strings.
Pydantic models	Every request body and response is a typed Pydantic model. No bare dicts.
DB sessions	Always use Depends(get_db). Never import db.engine directly in routes.
Idempotency	POST endpoints that create resources must be idempotent on duplicate IDs.

Frontend Development

The frontend is a React 18 SPA built with Vite 5 and Tailwind CSS 3. The UI proxies all `/api/` requests to `localhost:8000`. Component structure follows a `pages` `components` `hooks` hierarchy.

ui/src — component hierarchy

```
ui/src/
├── pages/
│   ├── HomePage.tsx           # Entry point – three mode tiles
│   ├── AnalyzePage.tsx       # Upload + recommendation display
│   ├── ShootModePage.tsx     # Step-by-step on-set instructions
│   ├── RecipesPage.tsx       # 13 recipe cards
│   ├── MyKitPage.tsx         # Equipment management
│   ├── SettingsPage.tsx      # User preferences
│   └── LabPage.tsx           # NGW Lab (restricted)
├── components/
│   ├── ui/                   # Primitive UI components
│   ├── analysis/             # Recommendation display cards
│   ├── shoot-mode/           # Step cards, role selector
│   └── lab/                   # Gold set, candidates, signals tabs
└── hooks/
    ├── useAnalysis.ts        # Upload + poll analysis result
    ├── useKit.ts             # My Kit CRUD state
    └── useLabAuth.ts         # Lab email auth state
```

bash — frontend dev commands

Start dev server with HMR

\$ cd ui && npm run dev

Type-check (no emit)

\$ cd ui && npm run typecheck

Lint

\$ cd ui && npm run lint

Production build (outputs to ui/dist/)

\$ cd ui && npm run build

Preview production build locally

\$ cd ui && npm run preview

Test Suite

The test suite uses pytest with a shared SQLite in-memory database fixture. Tests are organized by domain and must not require a running server or live VLM calls. VLM calls are mocked via pytest-mock.

tests/ — directory layout

```
tests/
■■■ conftest.py           # Shared fixtures: db, client, mock_vlm
■■■ test_engine.py        # Pattern detection + blueprint generation
■■■ test_shoot_match.py   # End-to-end shoot match service
■■■ test_signals.py       # Signal ingestion, hygiene flags
■■■ test_candidates.py    # Candidate CRUD and approval workflow
■■■ test_gold_sets.py     # Gold set CRUD and evaluation
■■■ test_reference_library.py # Reference library ingestion
■■■ test_perception_layer.py # Face validation, edge-case flags
■■■ test_learning.py      # Failure detection + severity scoring
```

bash — running tests

```
# Full suite (fastest — ~30 s)
$ python3 -m pytest tests/ -q --tb=short

# Single file with verbose output
$ python3 -m pytest tests/test_engine.py -v

# One specific test by name
$ python3 -m pytest tests/test_signals.py::test_hygiene_flags -v

# With coverage report
$ python3 -m pytest tests/ --cov=engine --cov=api --cov-report=term-missing

# Stop on first failure
$ python3 -m pytest tests/ -x
```

conftest.py — key fixtures

```
@pytest.fixture
def db():
    """In-memory SQLite database, created fresh per test."""
    engine = create_engine("sqlite:///memory:")
    Base.metadata.create_all(engine)
    with Session(engine) as session:
        yield session

@pytest.fixture
def mock_vlm(mockler):
    """Patch VLM gateway so tests never make real API calls."""
    return mockler.patch(
        "engine.vlm_adapter.gateway.analyze_image",
        return_value=MOCK_VLM_RESPONSE
    )
```

Benchmark System

The benchmark system runs the engine against a curated gold set and reports PASS / SOFT_PASS / FAIL per image. Results gate the CI/CD pipeline — any regression in gold set score below 75% fails the build.

bash — running benchmarks

```
# Standard benchmark run (all patterns)
$ python3 scripts/run_benchmarks.py

# Single pattern only
$ python3 scripts/run_benchmarks.py --pattern loop

# Gold sets only (skip live-data benchmarks)
$ python3 scripts/run_benchmarks.py --gold-sets-only

# JSON output for CI parsing
$ python3 scripts/nightly_benchmark.py --output=json > bench.json

# CI wrapper — exits 1 if any gold set score drops
$ bash scripts/ci_benchmark.sh
```

Result	Meaning	Action Required
PASS	Detected pattern matches gold label	None
SOFT_PAS S	Correct family, minor variant difference	Review if count increases
FAIL	Wrong pattern or confidence below threshold	Investigate + file candidate
NOT_TEST ED	Pattern has no gold set images yet	Add gold sets via Lab
ERROR	Engine raised an exception on this image	Fix before merging

Contribution Workflow

All contributions follow a branch-per-feature workflow with mandatory tests and benchmark validation before merge. No direct commits to main.

1

Branch

Create a feature branch: `git checkout -b feat/your-description`

2

Develop

Make changes. Keep commits focused. Run tests frequently with `pytest tests/ -q`

3

Test

All existing tests must pass. New features need new tests. Coverage should not drop.

4

Benchmark

Run `python3 scripts/run_benchmarks.py`. No new FAIL results allowed.

5

Validate YAMLS

If blueprints changed: `python3 scripts/validate_system_yamls.py --strict`

6

Pull Request

Open a PR against main. Include benchmark output and test summary in description.

Review

7

At least one curator review required. High-risk changes (solver, VLM) need two.

Merge

8

Squash-merge after approval. CI re-runs benchmarks as a gate before merge.

WHAT REQUIRES EXTRA REVIEW

- Changes to engine/solver.py — affects all pattern confidence scores.
- Changes to engine/blueprints/*.yaml — affects recommendation output.
- Changes to signal inclusion flags — affects learning and metrics eligibility.
- New VLM prompts — requires before/after benchmark comparison in PR.
- Any change that touches session_signals schema — requires migration script.

Environment Setup for New Developers

New developers need three things: a working Python environment, a populated database, and access to the AI Gateway.

bash — complete new developer setup

```
# 1. Clone and enter the repo
$ git clone https://github.com/your-org/ngw-core.git && cd ngw-core

# 2. Python environment
$ python3 -m venv .venv && source .venv/bin/activate
$ pip install -r requirements.txt

# 3. Frontend
$ cd ui && npm install && cd ..

# 4. Copy and fill in the env file
$ cp .env.example .env.local
# Edit .env.local — fill in AI_GATEWAY_API_KEY and your email

# 5. Initialize and seed the database
$ python3 -c "from db.database import init_db; init_db()"
$ python3 scripts/seed_starter_dataset.py

# 6. Verify everything works
$ python3 scripts/verify_dataset.py
$ python3 -m pytest tests/ -q
$ python3 scripts/run_benchmarks.py

# 7. Start development servers
$ uvicorn main:app --reload &
$ cd ui && npm run dev
```

GETTING LAB ACCESS

- Lab access requires your email in LAB_ALLOWED_EMAILS. Ask your team admin to add it and restart the backend.
- Then in the app: Settings Developer Tools enable, enter your email.
- The Lab tab appears immediately.